

# Operating Systems

## Scheduling

Course 211  
Spring Term 2018-2019

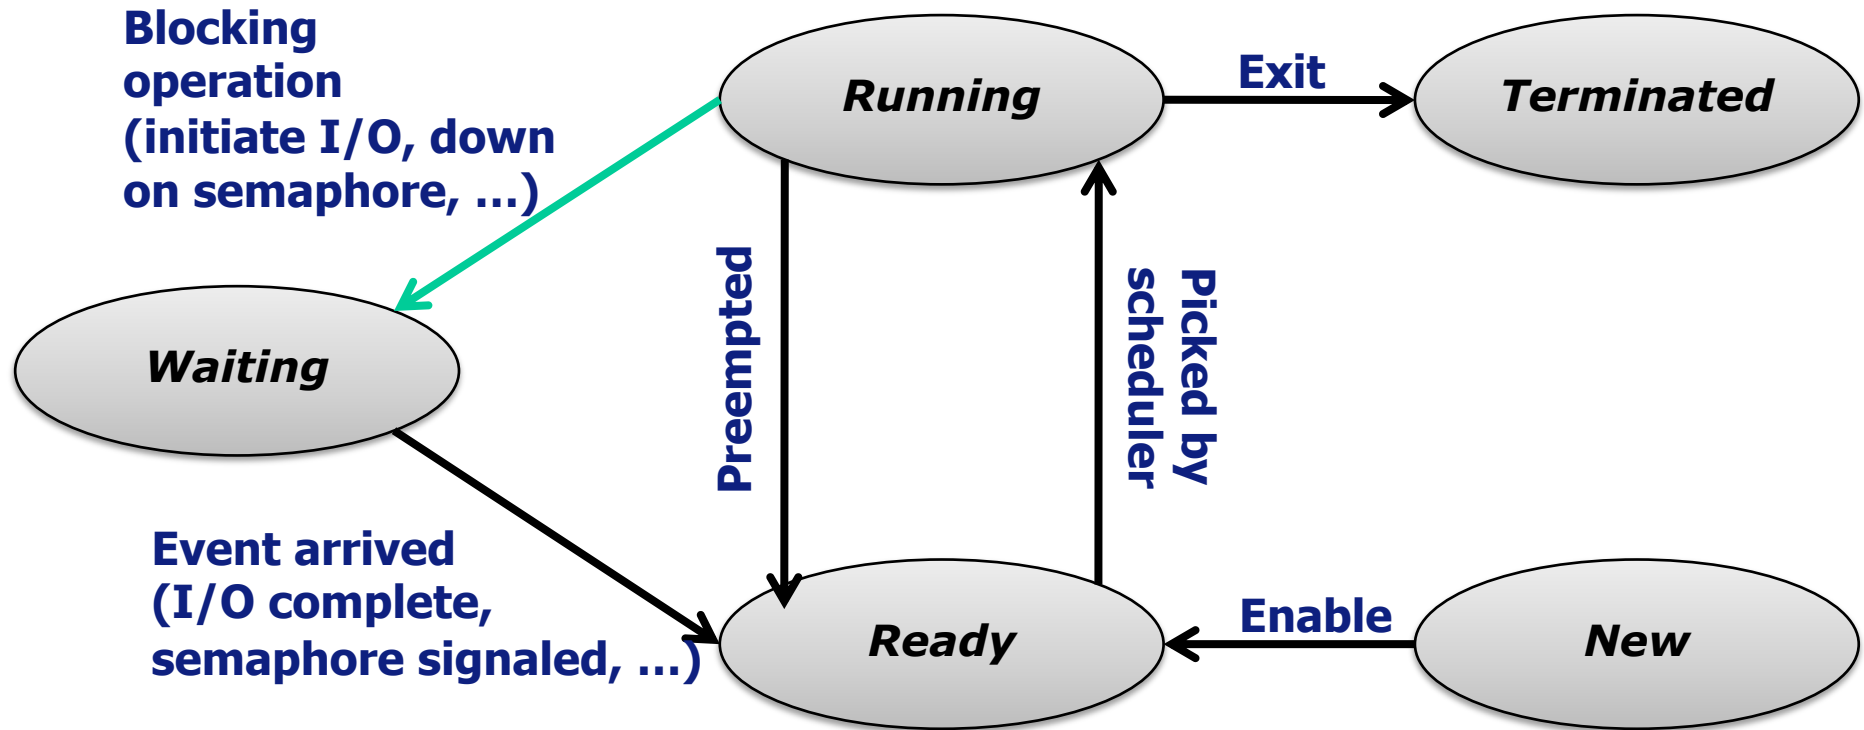
<http://www.imperial.ac.uk/computing/current-students/courses/211/calendar/>

(Slides courtesy of Cristian Cadar)

**Peter Pietzuch**

prp@doc.ic.ac.uk  
<http://www.doc.ic.ac.uk/~prp>

# Process States



- **New**: the process is being created
- **Ready**: runnable & waiting for processor
- **Running**: executing on a processor
- **Waiting/Blocked**: waiting for an event
- **Terminated**: process is being deleted

**If multiple processes are ready, which one should be run?**

# Goals of Scheduling Algorithms

## Ensure fairness

- Comparable processes should get comparable services

## Avoid indefinite postponement

- No process should starve

## Enforce policy

- E.g. priorities

## Maximize resource utilisation

- CPU, I/O devices

## Minimize overhead

- From context switches, scheduling decisions

# Goals of Scheduling Algorithms

## **Batch systems:**

- Throughput: Jobs per unit of time
- Turnaround time: Time between job submission & completion

## **Interactive systems:**

- Response time crucial: Time between request issued and first response

## **Real-time systems:**

- Meeting deadlines:
  - Soft deadlines: e.g. leads to degraded video quality
  - Hard deadline: e.g. leads to plane crash

# Preemptive vs. Non-Preemptive Scheduling

## **Non-preemptive**

- Let process run until it blocks or voluntarily releases CPU

## **Preemptive**

- Let process run for a maximum amount of fixed time
- Requires clock interrupt

# CPU-Bound vs. I/O-Bound Processes

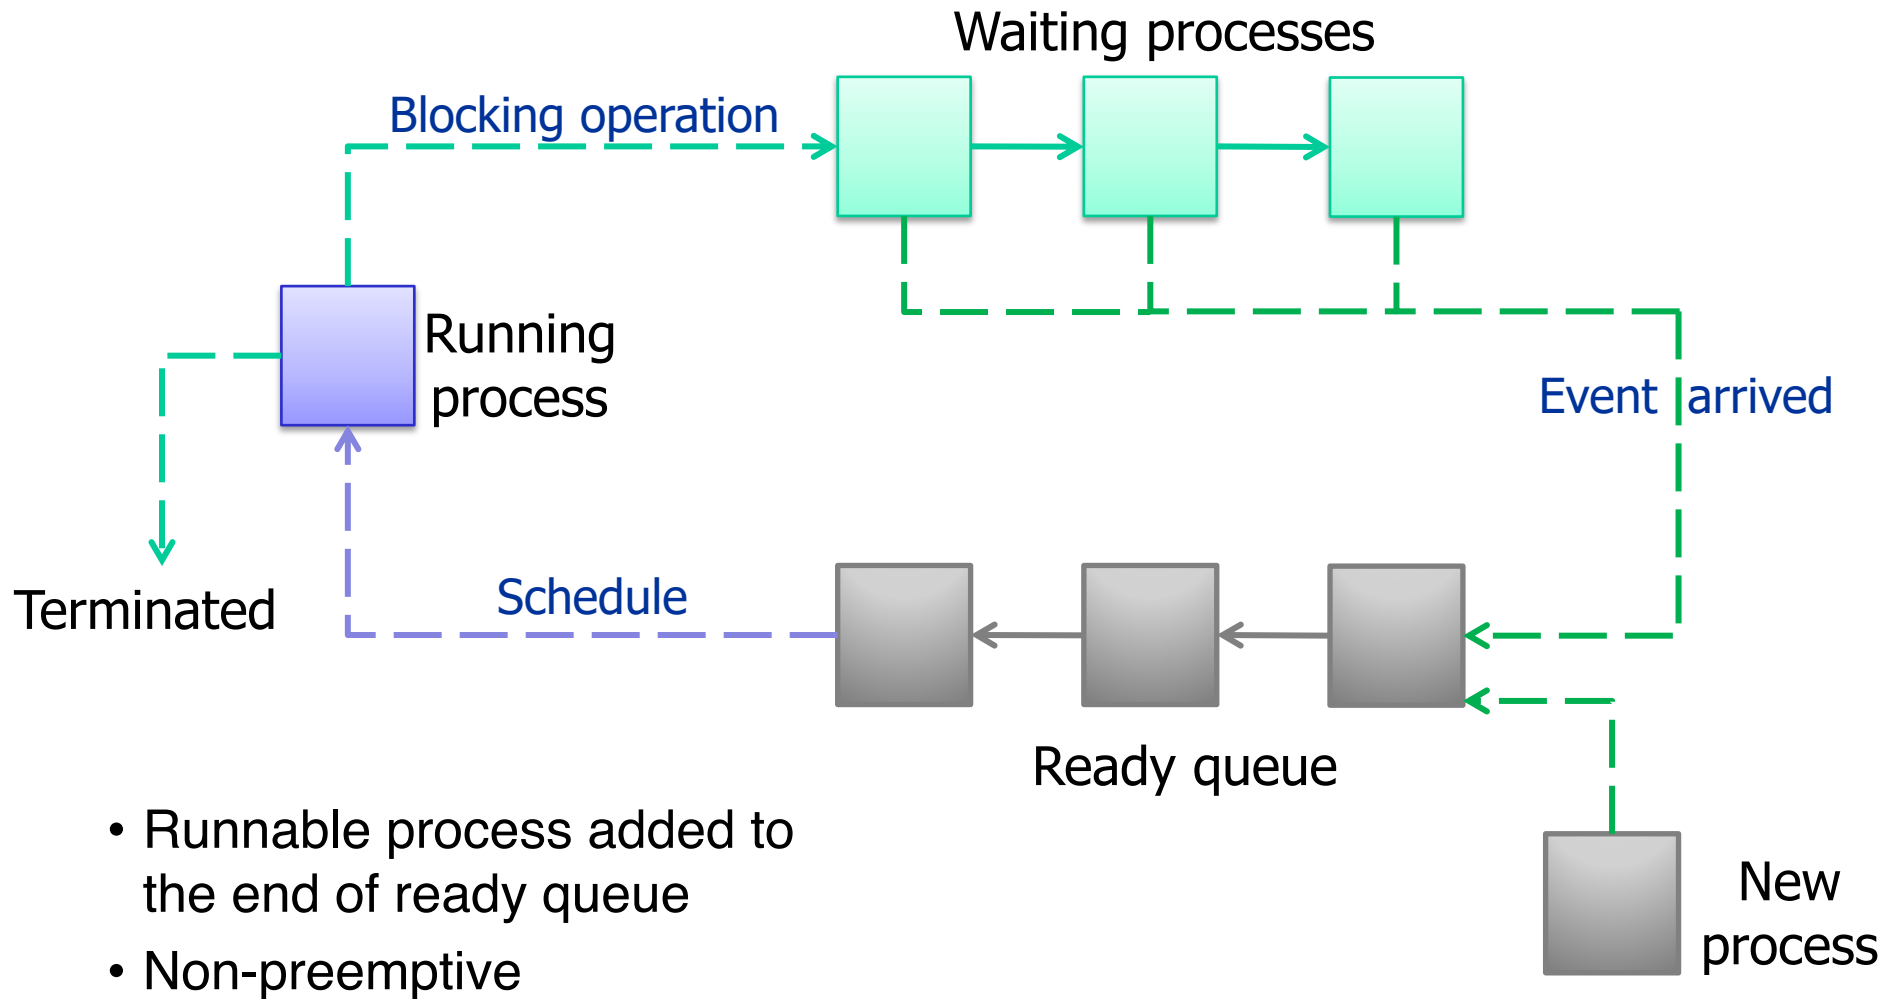
## **CPU-bound processes**

- Spend most of their time using the CPU

## **I/O-bound processes**

- Spend most of their time waiting for I/O
- Tend to only use CPU briefly before issuing I/O request

# First-Come First-Served (FCFS) (Non-Preemptive)



# FCFS Advantages

No indefinite postponement

- All processes are eventually scheduled


Really easy to implement




# FCFS Disadvantages

What happens if a long job is followed by many short jobs?

- E.g. 1h, 1s, 1s, 1s, with jobs 2-4 submitted just after job 1

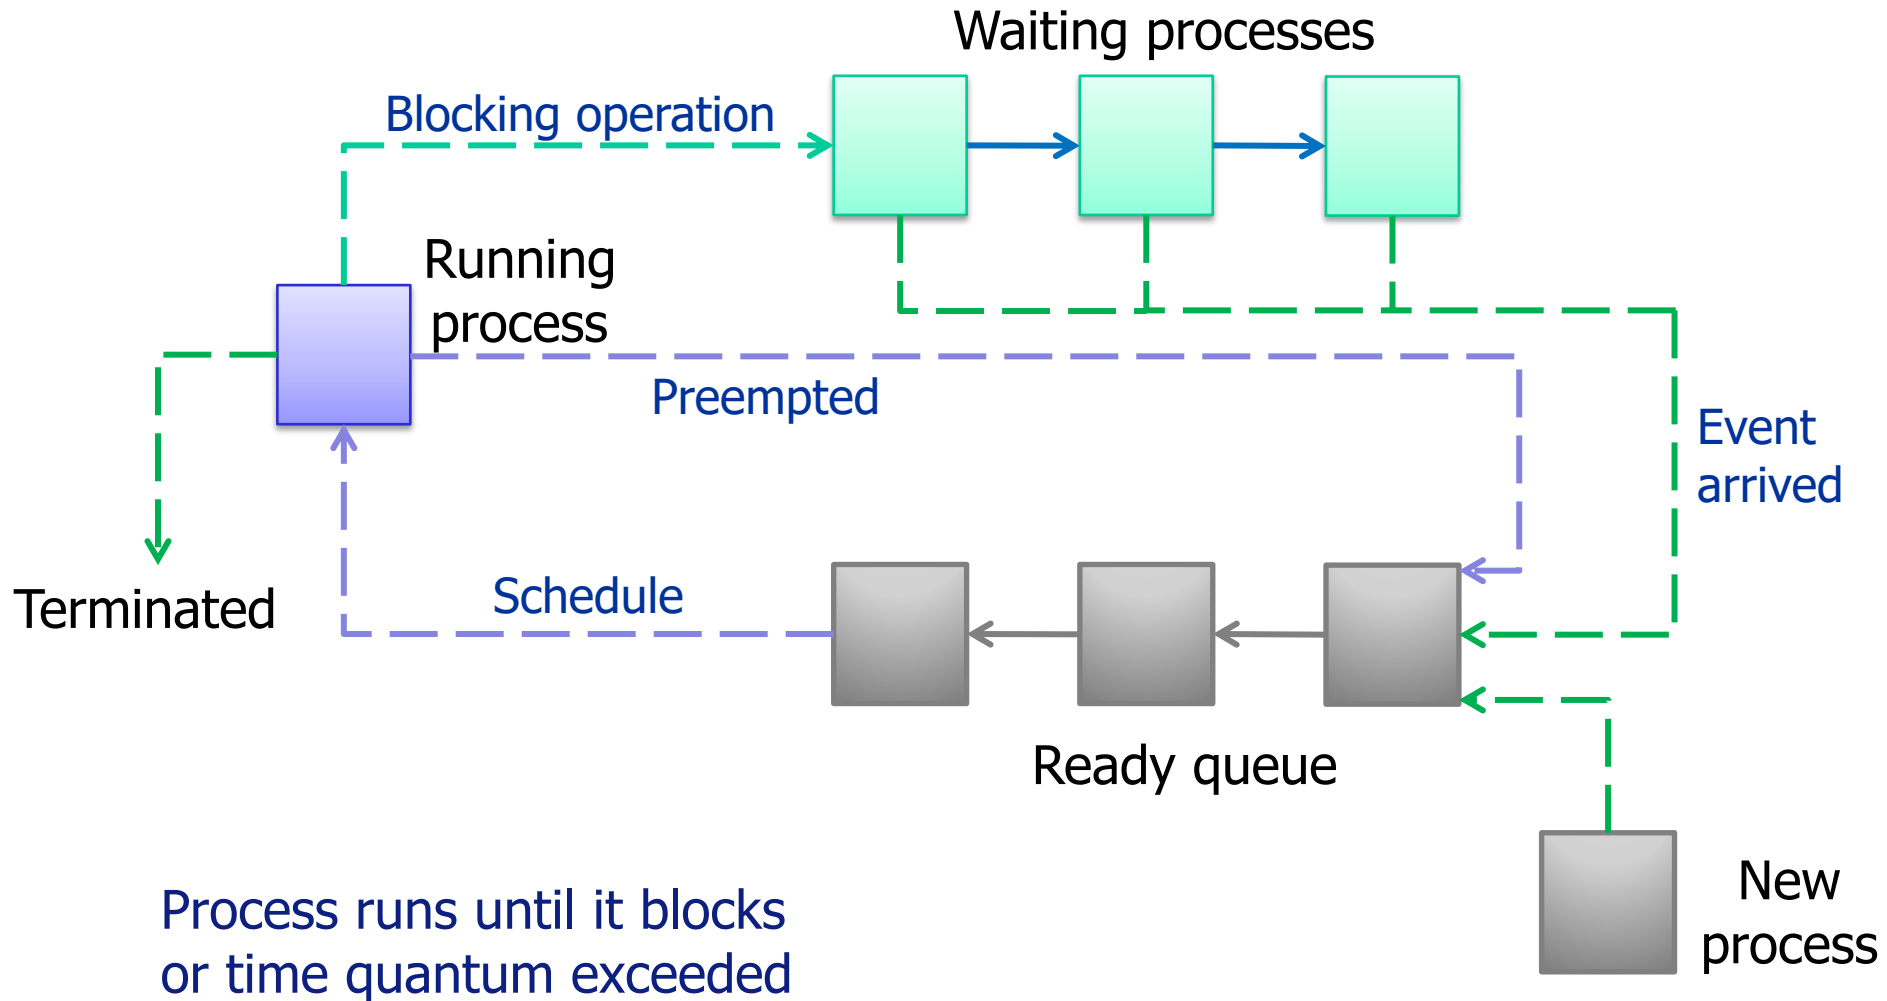


- Throughput?
- Average turnaround time?



- Throughput?
- Average turnaround time?

# Round-Robin Scheduling (RR)



# Round-Robin

## Fairness

- Ready jobs get equal share of the CPU

## Response time

- Good for small number of jobs

## Average turnaround time:

- Low when run-times differ
- Poor for similar run-times

# RR Quantum (Time Slice)

## RR Overhead:

- 4 ms quantum, 1 ms context switch time:  
20% of time  $\rightarrow$  overhead
- 1 s quantum, 1ms context switch time:  
only 0.1% of time  $\rightarrow$  overhead

## Large quantum:

- Smaller overhead
- Worse response time
  - Quantum =  $\infty \rightarrow$  FCFS

## Small quantum:

- Larger overhead
- Better response time

# RR Quantum

## Choosing a quantum value:

- Should be much larger than context switch cost
- But provide decent response time

Typical values: 10 ms – 200 ms

Some example values for standard processes (values vary depending on process type and behaviour, priority, ...):


- Linux: 100 ms
- Windows client: 20 ms
- Windows server: 180 ms

# Shortest Job First (SJF)


Non-preemptive scheduling with run-times known in advance

Pick the shortest job first

- Process with shortest CPU burst



FCFS turnaround time:



SJF turnaround time:

- Provably optimal when all jobs are available simultaneously

# Shortest Remaining Time (SRT)

Preemptive version of shortest job first


- Again, runtimes have to be known in advance

Choose process whose remaining time is shortest

- When new process arrives with execution time less than the remaining time for the running process, run it

Allows new short jobs to get good service

Example: 3 jobs:  $J_1 = 6s$ ,  $J_2 = 4s$ ,  
 $J_3 = 1s$  arrives after  $2s$



# Shortest Remaining Time (SRT)

What if a running process is almost complete and a shorter job arrives?



# Knowing Run-times in Advance

Run-times are usually not available in advance

Compute CPU burst estimates based on heuristics?

- E.g. based on previous history
- Not always applicable

User-supplied estimates?

- Need to counteract cheating to get higher priority
- E.g. terminate or penalise processes after they exceed their estimated run-time

# Fair-Share Scheduling

Users are assigned some fraction of the CPU

- Scheduler takes into account who owns a process before scheduling it

E.g. two users each with 50% CPU share

- User 1 has 4 processes: A, B, C, D
- User 2 has 2 processes: E, F

What does a fair-share RR scheduler do?

# Priority Scheduling

Jobs are run based on their priority

- Always run the job with the highest priority

Priorities can be **externally defined** (e.g. by user) or based on some **process-specific metrics** (e.g. their expected CPU burst)

Priorities can be **static** (i.e. they do not change) or **dynamic** (they may change during execution)

Example: Consider 3 processes arriving at essentially the same time with externally defined static priorities  $A = 4$ ,  $B = 7$ ,  $C = 1$ , where a higher value means higher priority.

- Processes are run to completion in the order B, A, C

# General-Purpose Scheduling

## Favour short and I/O-bound jobs

- Get good resource utilisation
- And short response times

## Quickly determine the nature of job and adapt to changes

- Processes have periods when they are I/O-bound and periods when they are CPU-bound

# Multilevel Feedback Queues

A form of priority scheduling

- Shortest remaining time also a form of priority scheduling

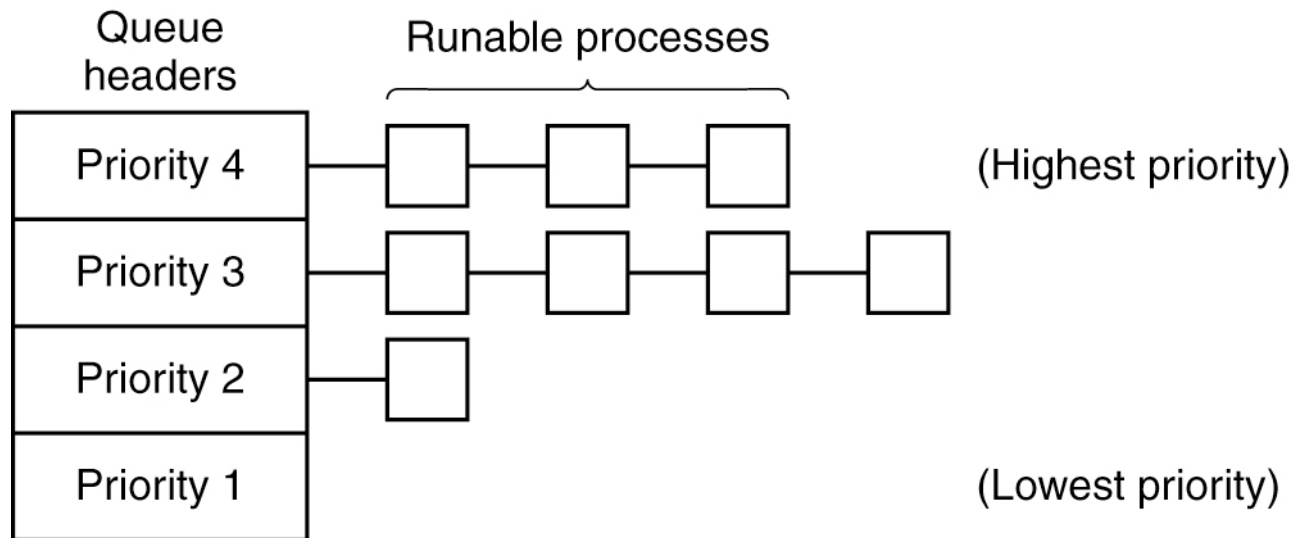
Implemented by many OSs:

- Windows Vista, Windows 7
- Mac OS X
- Linux 2.6 - 2.6.23
- Pintos!

# Multilevel Feedback Queues

## One queue for each priority level

- Run job on highest non-empty priority queue
- Each queue can use different scheduling algorithm
  - Usually round-robin



# Multilevel Feedback Queues

Need to determine current nature of job

- I/O-bound? CPU-bound?

Need to worry about starvation of lower-priority jobs

Feedback mechanism:

- Job priorities recomputed periodically, e.g. based on how much CPU they have recently used
  - Exponentially-weighted moving average
- **Aging**: increase job's priority as it waits

# Multilevel Feedback Queues

## Not very flexible

- Applications basically have no control
- Priorities make no guarantees
  - What does priority 15 mean?

## Does not react quickly to changes

- Often needs warm-up period
  - Running system for a while to get better results
- Problem for real-time systems, multimedia apps

## Cheating is a concern

- Add meaningless I/O to boost priority?

## Cannot donate priority



# Lottery Scheduling [Waldspurger and Weihl 1994]

Jobs receive lottery tickets for various resources

- E.g. CPU time

At each scheduling decision, one ticket is chosen at random and the job holding that ticket wins

Example: 100 lottery tickets for CPU time,  
P1 has 20 tickets

- Chance of P1 running during the next CPU quantum: 20%
- In the long run, P1 gets 20% of the CPU time

# Lottery Scheduling

## Number of lottery tickets meaningful

- Job holding  $p\%$  of tickets, gets  $p\%$  of resource
- Unlike priorities

## Highly responsive:

- New job given  $p\%$  of tickets has the  $p\%$  chance to get the resource at the **next** scheduling decision

## No starvation

## Jobs can exchange tickets

- Allows for priority donation
- Allows cooperating jobs to achieve certain goals

## Adding/removing jobs affect remaining jobs proportionally

## Unpredictable response time

- What if interactive process is unlucky for a few lotteries?

# Summary

Scheduling algorithms often need to balance conflicting goals

- E.g. ensure fairness, enforce policy, maximise resource utilisation

Different scheduling algorithms appropriate in different contexts

- E.g. batch systems vs interactive systems vs real-time systems

Well-studied scheduling algorithms include:

- First-Come First-Served FCFS, Round Robin, Shortest Job First (SJF), Shortest Remaining Time (SRT), Multilevel Feedback Queues and Lottery Scheduling

# Tutorial Questions

State which of the following are true and which false. Justify your answers.

1. Interactive systems generally use non-preemptive processor scheduling.
2. Turnaround times are more predictable in preemptive than in non-preemptive systems.
3. One weakness of priority scheduling is that the system will faithfully honor the priorities, but the priorities themselves may not be meaningful.

# Tutorial Questions

1. Interactive systems generally use non-preemptive processor scheduling.
2. Turnaround times are more predictable in preemptive than in non-preemptive systems.
3. One weakness of priority scheduling is that, while a system may faithfully honour the priorities, the priorities themselves may not be meaningful.