

Operating Systems Device Management



Course 211
Spring Term 2018-2019

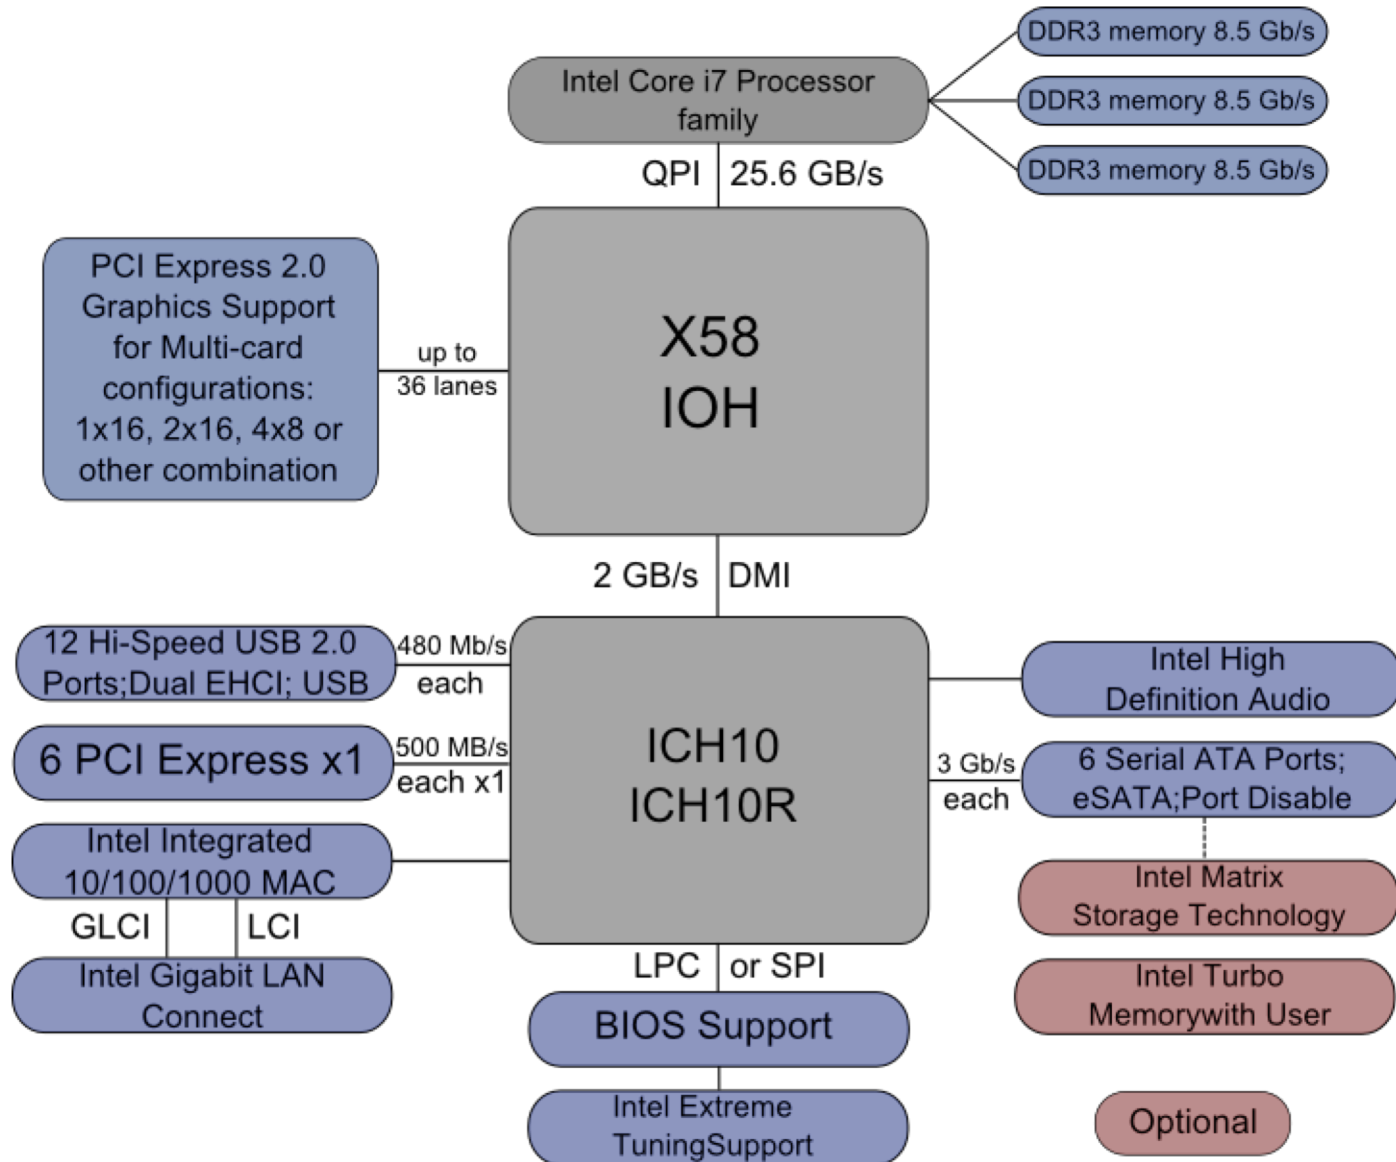
<http://www.imperial.ac.uk/computing/current-students/courses/211/calendar/>

Based on slides by Daniel Rueckert

Peter Pietzuch

prp@doc.ic.ac.uk
<http://www.doc.ic.ac.uk/~prp>

Example: Intel Architecture



I/O Device Management

Objectives

- Fair access to shared devices
 - Allocation of dedicated devices
- Exploit parallelism of I/O devices for multiprogramming
- Provide uniform simple view of I/O
 - Hide complexity of device handling
 - Give uniform naming and error handling

Device Independence


Device independence from

- Device type (e.g. terminal, disk or DVD drive)
- Device instance (e.g. which disk)

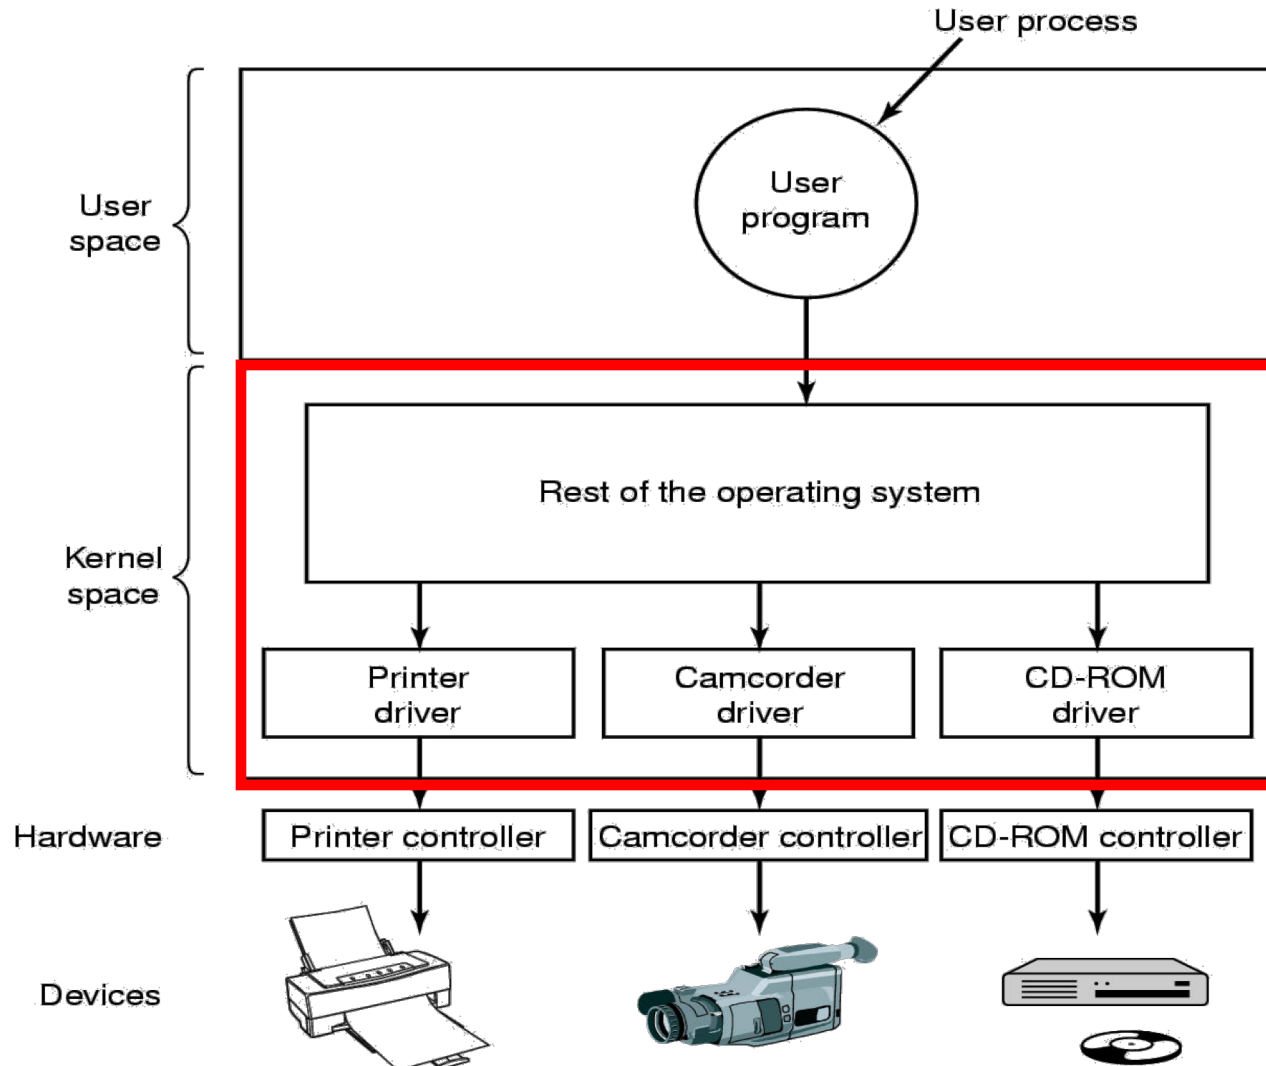
Device variations

- Unit of data transfer: **character** or **block**
- Supported operations: e.g. **read**, **write**, **seek**
- **Synchronous** or **asynchronous** operation
- Speed differences
- Sharable (e.g. disks) or single user (e.g. printer, DVD-RW)
- Types of error conditions

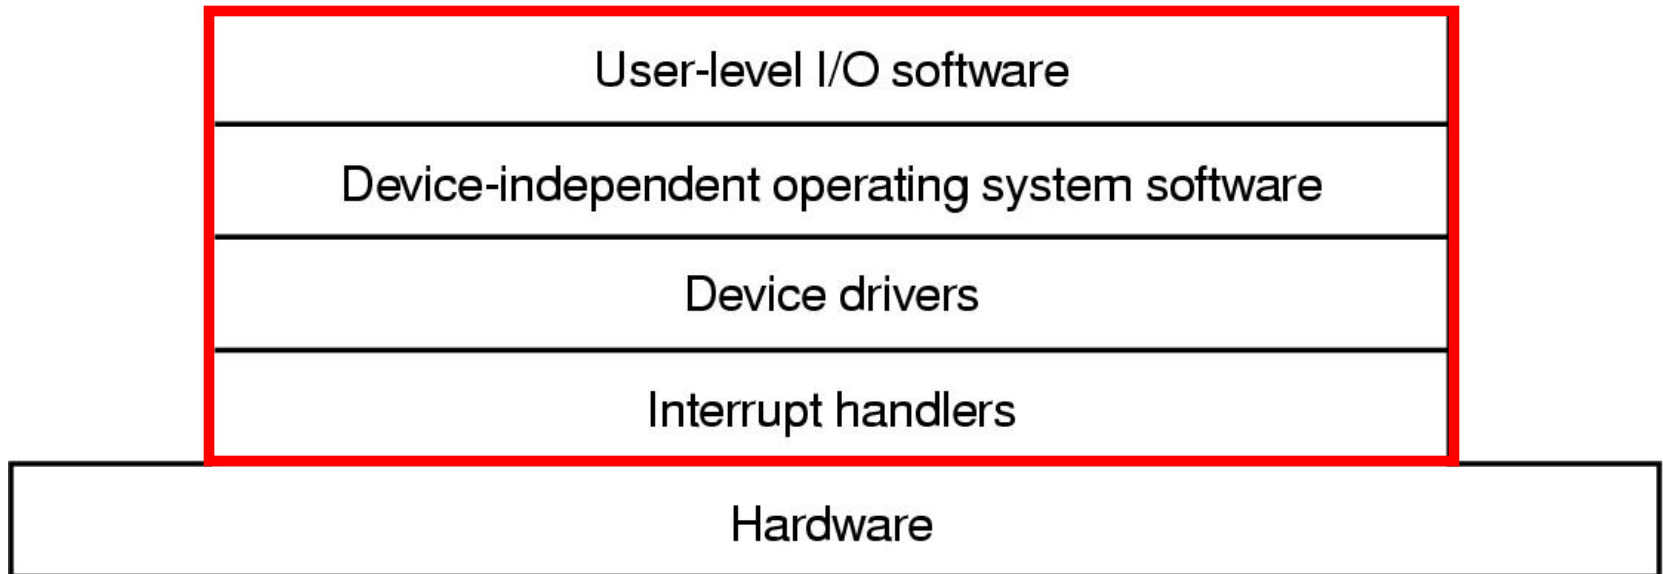
Device Variations: Character vs. Block



I/O Layering



I/O Layers: Overview



Interrupt Handler

Interrupt handler

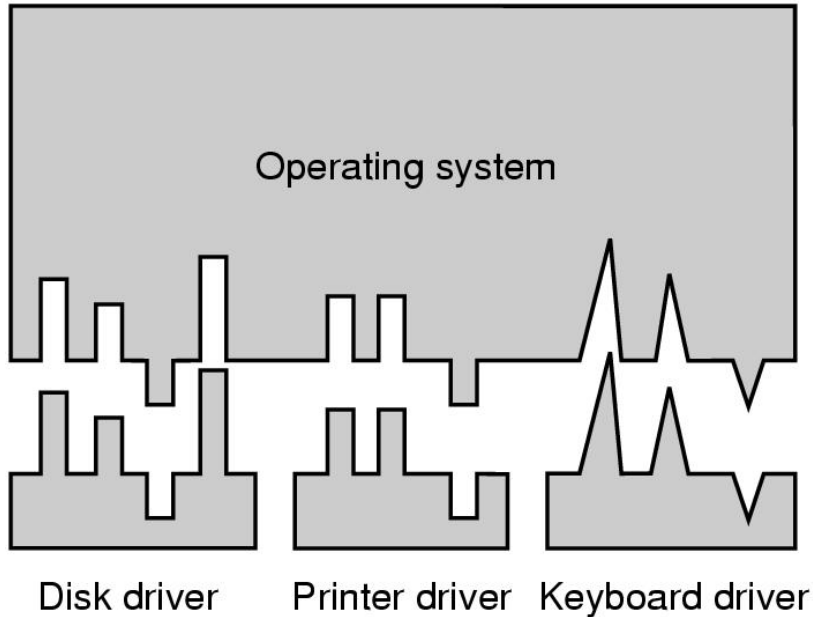
- Process each interrupt
- For **block** devices:
 - on transfer completion, signal device handler
- For **character** devices
 - when character transferred, process next character

Device Driver

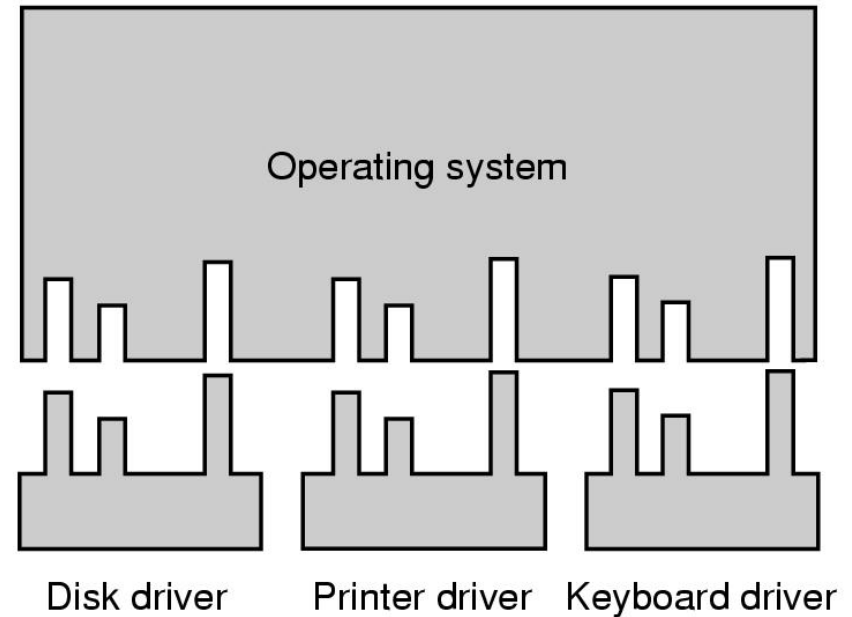
Device handler/driver

- Handles one device type
 - but may control multiple devices of same type
- Implements block read or write
- Access device registers
- Initiate operations
- Schedule requests
- Handle errors

Device Independent OS Layer I



(a)



(b)

(a) without standard driver interface

(b) with standard driver interface

Device Independent OS Layer II

Device independent layer provides device independence

- Mapping logical to physical devices (naming and switching)
- Request validation against device characteristics
- Allocation of dedicated devices
- Protection/user access validation
- Buffering for performance and block size independence
- Error reporting

Dedicated vs. Shared Device Allocation

Dedicated device (e.g. DVD writer, terminal, printer, ...)

- Simple policy:
 - Open fails if already opened
 - Alternatively, queue open requests
- Allocated for long periods
- Only allocated to authorised processes

Shared device (e.g. disks, window terminals, ...)

- OS provides file system for disks

Device Allocation: Spooling

Blocking user access to allocated, nonsharable devices?

- Causes delays and bottlenecks

- ☛ Spool to intermediate medium (disk file)

Spooled devices (e.g. printers)

1. Printer output saved to disk file


2. File printed later by **spooler daemon**

- Printer only allocated to spooler daemon
- No normal process allowed direct access

- Provides sharing of nonsharable devices

- Reduces I/O time ➔ gives greater throughput

Buffering



Buffered vs. Unbuffered I/O

Buffered I/O

Output: User data transferred to OS output buffer

Process continues and only suspends when buffer full

Input: OS reads ahead; reads normally satisfied from buffer

Process blocks when buffer empty

- Used to smooth peaks in I/O traffic
- Caters for differences in data transfer units between devices

Unbuffered I/O

- Data transferred directly from user space to/from device
 - Each read/write causes physical I/O
 - Implies device handler used for each transfer
- High process switching overhead (e.g. per character)

User-Level I/O Interface

User interface

- I/O operations: `open`, `close`, `read`, `write`, `seek`
- OS I/O library procedures to set up parameters
 - Must be device independent
- Synchronous or asynchronous
- Blocking or non-blocking
- Unix: Access virtual devices as files

Device Drivers


Memory-Mapped I/O

Device addressed as memory location

Example: Disabling the I2S clock on Raspberry PI:

```
* (clk+0x26) = 0x5A000000;
```

```
* (clk+0x27) = 0x5A000000;
```



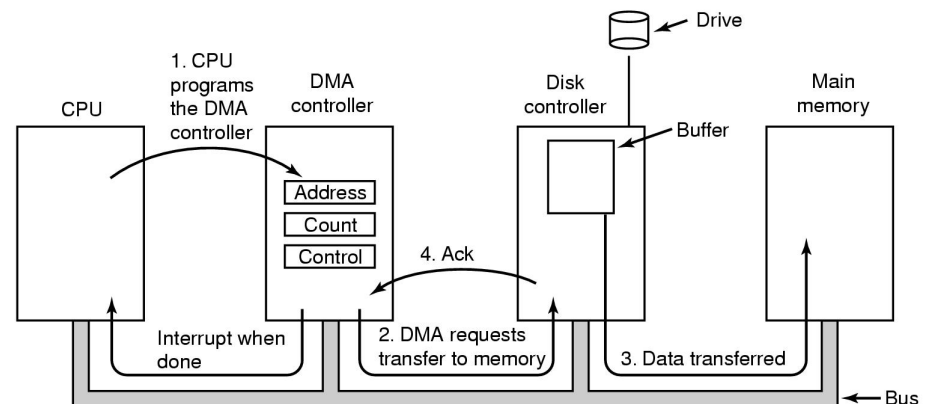
More flexible

Ways to do I/O

1. Programmed I/O

2. Interrupt-Driven I/O

3. I/O using Direct Memory Access (DMA)



Linux: Loadable Kernel Module (LKM)

Loadable kernel modules provide device drivers

- Contain object code, loaded **on-demand**
 - Dynamically linked to running kernel
 - Provided by hardware vendors or independent developers
- Require **binary compatibility**
 - Modules written for different kernel versions may not work

Kmod

- Kernel subsystem managing modules without user intervention
- Determines module dependencies
- Load modules on demand

Linux: Basic LKM module

Every LKM consists of two basic functions (minimum):

```
int init_module(void) /* used for all initialisation code */
{
...
}
void cleanup_module(void) /* used for clean shutdown */
{
...
}
```

Load module by issuing following command:

```
insmod module.o
```

- Normally restricted to root

Linux I/O Management

Linux I/O Management

Kernel provides common interface for I/O system calls

Devices grouped into **device classes**

- Members of each device class perform similar functions
- Allows kernel to address performance needs of certain devices (or classes of devices) individually

Major and **minor** identification numbers

- Used by device drivers to identify their devices
- Devices with same **major** num controlled by same driver
- **Minor** nums enable system to distinguish between devices of same class

Linux: Device Drivers

Device special files

- Most devices represented by device special files
- Entries in `/dev` directory that provide access to devices
- List of devices in system can be obtained by reading contents of `/proc/devices`:

Character devices:

```
1 mem
2 pty
4 ttyS
5 cua
10 misc
13 input
109 lvm
136 pts
162 raw
180 usb
```

Block devices:

```
1 ramdisk
2 fd
3 ide0
7 loop
8 sd
9 md
58 lvm
65 sd
66 sd
```

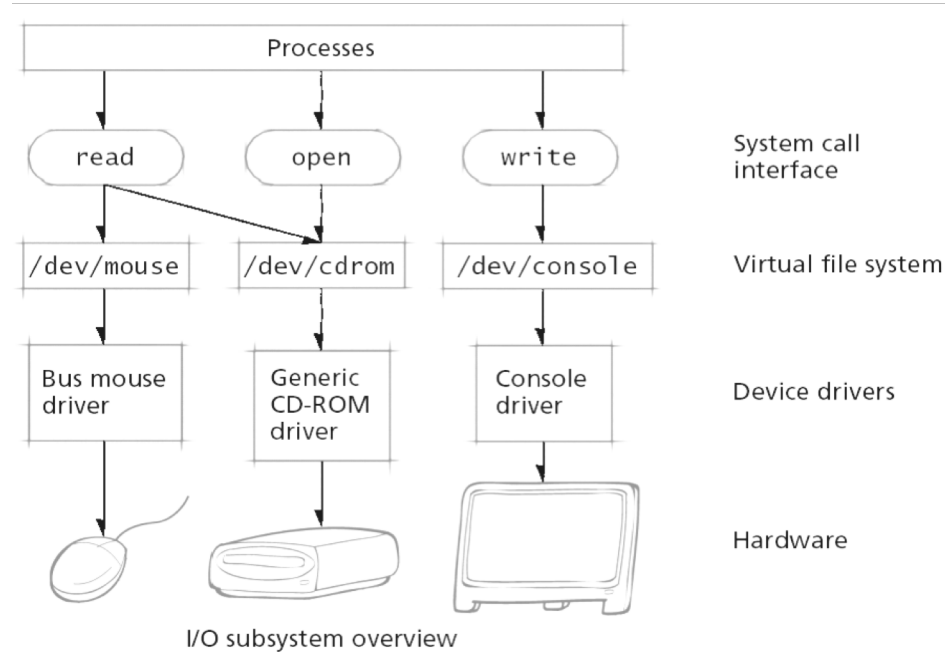

Linux: /dev

c/b			major minor		file name		
↓			↓	↓		↓	
crw-----	1	root	root	5,	1	Dec 27 16:09	console
brw-rw-rw-	1	root	disk	2,	0	May 21 2001	fd0
brw-rw-rw-	1	root	disk	2,	4	May 21 2001	fd0d360
brw-rw-rw-	1	root	disk	2,	8	May 21 2001	fd0h1200
brw-rw-rw-	1	root	disk	2,	40	May 21 2001	fd0h1440
crw-rw----	1	root	lp	6,	0	May 21 2001	lp0
crw-rw----	1	root	lp	6,	1	May 21 2001	lp1
crw-rw----	1	root	lp	6,	2	May 21 2001	lp2
crw-rw----	1	root	lp	180,	0	May 21 2001	usb1p0
crw-rw----	1	root	lp	180,	1	May 21 2001	usb1p1
crw-rw----	1	root	lp	180,	2	May 21 2001	usb1p2
lrwxrwxrwx	1	root	root		10	Dec 6 06:53	mouse -> /dev/psaux
crw-rw-r--	1	root	root	10,	1	May 21 2001	psaux
lrwxrwxrwx	1	root	root		3	Nov 30 2001	cdrom -> hdc
brw-rw-rw-	1	root	disk	3,	0	May 21 2001	hda
brw-rw-rw-	1	root	disk	3,	16	May 21 2001	hdb
brw-rw-rw-	1	root	disk	3,	32	May 21 2001	hdc

Linux: Device Access

Device files accessed via **virtual file system (VFS)**

- System calls pass to VFS, which in turn issues calls to device drivers
- Most drivers implement common file operations
 - e.g. read, write, seek



Linux provides **ioctl** system call

- Supports special tasks:
 - Ejecting CD-ROM tray
`ioctl(cdrom, CDROMEJECT, 0)`
 - Retrieving status information from printer

Linux: Character Device I/O I

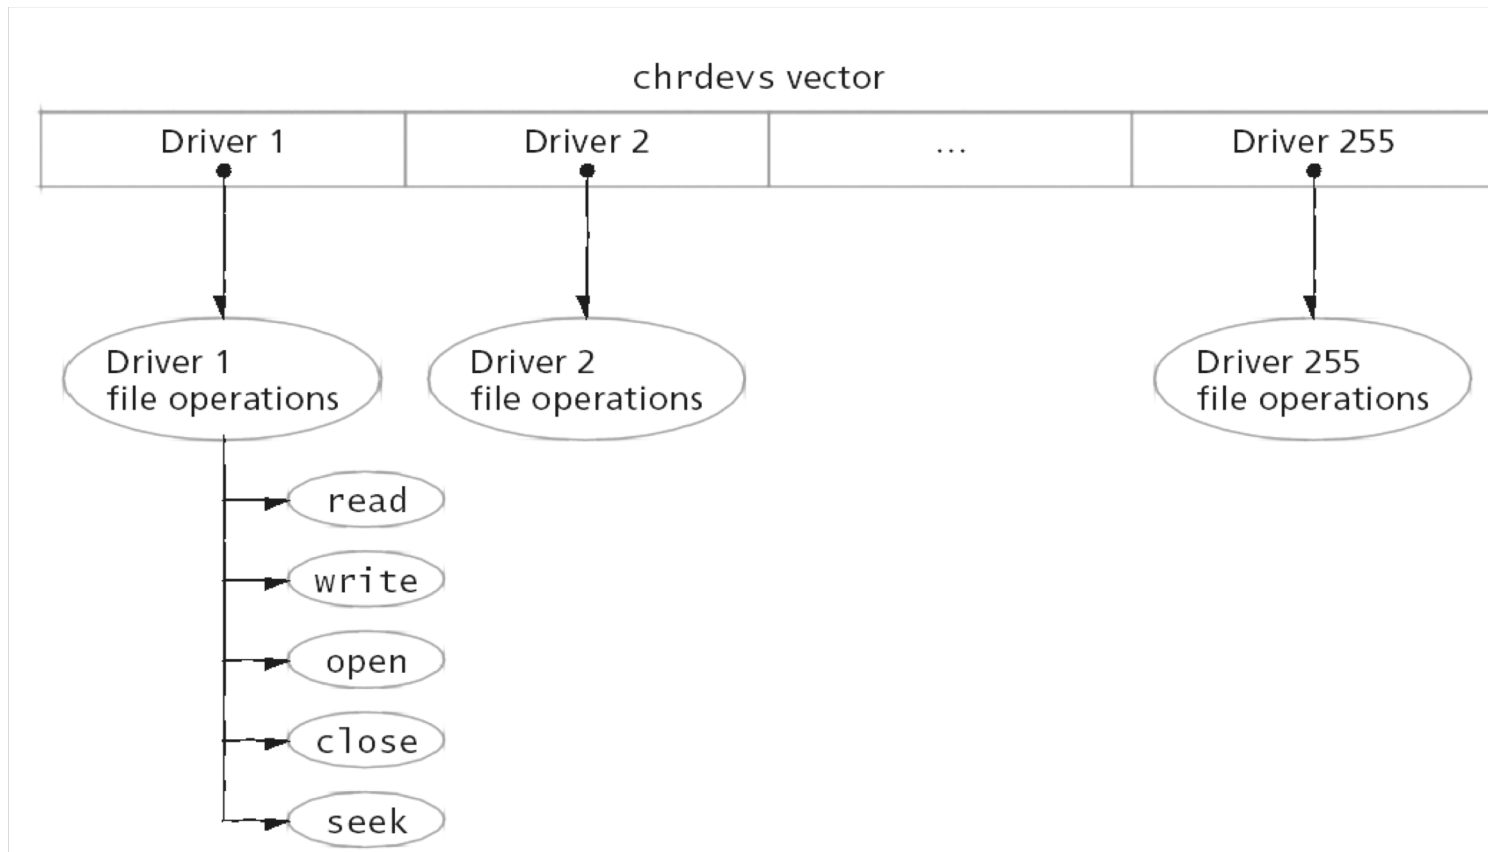
Character device

- Transmits data as **stream of bytes**
- Represented by `device_struct` structure contains:
 - Driver name
 - Pointer to driver's `file_operations` structure
- All registered drivers referenced by `chrdevs` vector

`file_operations` structure

- Maintains operations supported by device driver
- Stores functions called by VFS when system call accesses device special file

Linux: Character Device I/O II



Linux: Block Device I/O

Block I/O subsystem

- Kernel's block I/O subsystem contains number of layers
- Modularise block I/O operations by placing common code in each layer

Two primary strategies used by kernel to minimise amount of time spent accessing block devices:

- Caching data
- Clustering I/O operations

Linux: Block Device Caching

When data from block device requested, kernel first searches **cache**

- If found, data copied to process's address space
- Otherwise, typically added to request queue

Direct I/O

- Driver bypasses kernel cache when accessing device
- Important for databases and other applications
 - Kernel caching inappropriate and may reduce performance/consistency

Linux I/O API

Linux I/O Classes

I/O classes

Character (unstructured):	Files and devices
Block (structured):	Devices
Pipes (message):	Interprocess communication
Socket (message):	Network interface

Sockets

Allow **bidirectional** communication

Can be used to exchange information both locally and across a network

- Unlike pipes which are identified by machine specific file descriptors

Two types of sockets:

- **TCP** (stream sockets)
- **UDP** (datagram sockets)

Linux I/O API I

I/O calls

```
fd = create(filename, permission)
```

Opens file for reading/writing; fd is index to file descriptor, permission is used for access control

```
fd = open(filename, mode)
```

Mode is 0, 1, 2 for read, write, read/write

Linux: I/O API II

close (fd)

Close file or device

numbytesread = read(fd, buffer, numbytes)

read **numbytes** from file or device referenced by **fd** into memory buffer; returns number of bytes actually read in **numbytesread**

numbyteswritten = write(fd, buffer, numbytes)

write **numbytes** to file referenced by **fd** from memory buffer; returns number of bytes actually written in **numbyteswritten**

Linux: I/O User Interface API III

`pipe (&fd[0])`

Creates pipe; `fd` is an array of two integers: `fd[0]` is for reading, `fd[1]` for writing

`newfd = dup (oldfd) , dup2 (oldfd, newfd)`

Duplicate file descriptor

`ioctl (fd, operation, &termios)`

Used to control devices; e.g. `&termios` is array of control chars

`fd = mknod (filename, permission, dev)`

Creates new special file e.g. character or block device

Linux: File Descriptors

Each process has its own **file descriptor table**

- Each process has 3 file descriptors when created:

file descriptor	input/output
0	stdin
1	stdout
2	stderr

- By default, all three file descriptors refer to terminal from which program was started

Linux: I/O Example I

```
#include <stdlib.h>
#define BUFSIZE 512

int main( int argc, char ** argv){

    int fd, n, stdin, stdout, stderr;
    char buffer[BUFSIZE];

    /* Standard input always corresponds to fd = 0 */
    stdin = 0;

    /* Standard output always corresponds to fd = 1 */
    stdout = 1;

    /* Standard error always corresponds to fd = 2 */
    stderr = 2;

    /* Open file */
    fd = open(argv[1], O_RDONLY);
```

Linux: I/O Example II

```
if (fd < 0) {  
    write(stderr, "Can't open file", 15);  
} else {  
    do {  
        n = read(fd, buffer, BUFSIZE);  
        if (n < 0) {  
            write(stderr, "Error while reading", 19);  
        } else {  
            write(stdout, buffer, n);  
        }  
    } while (n > 0);  
}  
  
/* Close file */  
close(fd);  
}
```

Blocking vs. Non-blocking I/O

Blocking I/O

- I/O call returns when operation completed
- Process suspended → I/O appears “instantaneous”
- Easy to understand but leads to multi-threaded code

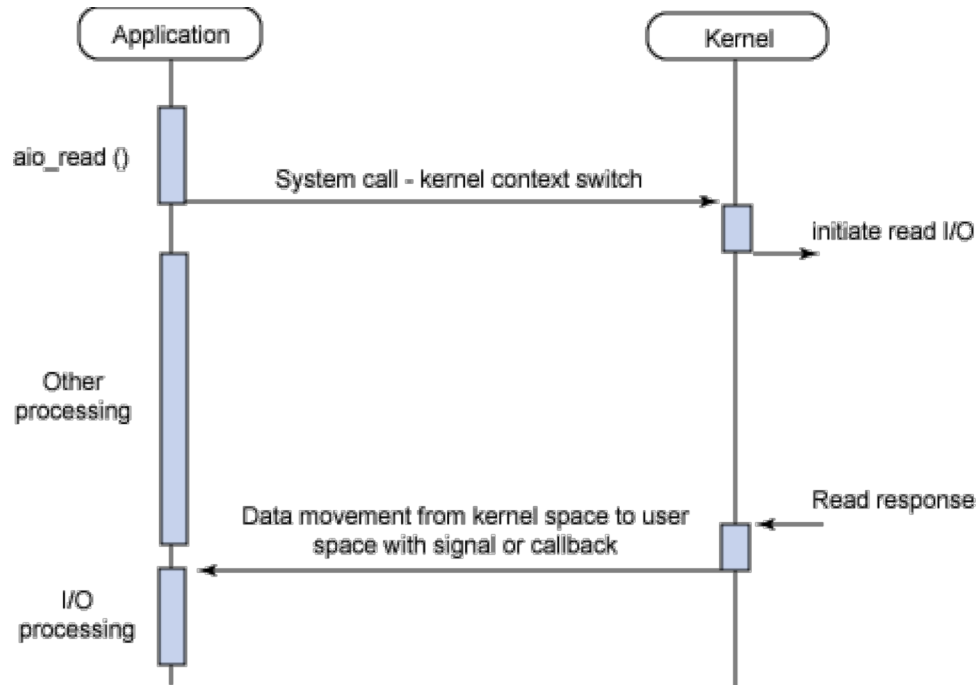
Non-blocking I/O

- I/O call returns as much as available (e.g. `read` with 0 bytes)
- Turn on for file descriptor using `fcntl` system call
- Provides application-level polling for I/O (**how?**)

Asynchronous I/O

Asynchronous I/O

- Process executes in parallel with I/O operation
 - No blocking in interface procedure
- I/O subsystems notifies process upon completion
 - Callback function, process signal, ...
- Supports check/wait if I/O operation completed
- Very flexible and efficient
- Harder to use and potentially less secure (why?)



Source: IBM Developerworks

Linux: AIO Example I

AIO: Support for asynchronous I/O in Linux 2.6

```
#include <aio.h>

...

int fd, ret;
struct aiocb my_aiocb;

fd = open("myfile", O_RDONLY );

/* Allocate buffer for aio request */
my_aiocb.aio_buf = malloc(BUFSIZE + 1);

/* Initialise aio control structure */
my_aiocb.aio_fildes = fd;
my_aiocb.aio_nbytes = BUFSIZE;
my_aiocb.aio_offset = 0;
```

Linux: AIO Example II

/* Initiate read request */

```
ret = aio_read(&my_aioCB);
```

/* Wait for read to finish (more usefully do something else)

Also possible to register signal notification or thread callback */

```
while (aio_error(&my_aioCB) == EINPROGRESS);
```

/* Check result from read */

```
if ((ret = aio_return(&my_aioCB)) > 0) {
```

/* Successfully read ret bytes */

```
} else {
```

/* Read failed, check errno*/

```
}
```

Tutorial Question

In which of the four I/O software layers (user-level I/O software, device-independent OS software, device drivers and interrupt handlers) is each of the following done?

- (a) Computing the track, sector and head for a disk read
- (b) Maintaining a cache of recently used blocks
- (c) Writing commands to the drive registers
- (d) Checking to see if the user is permitted to use the device
- (e) Converting binary integers to ASCII for printing

Tutorial Answer

- (a) Computing the track, sector and head for a disk read
- (b) Maintaining a cache of recently used blocks
- (c) Writing commands to the drive registers
- (d) Checking to see if the user is permitted to use the device
- (e) Converting binary integers to ASCII for printing